

<https://github.com/EdouardDabo/DIA>

Telco Customer churn dataset

Code explanation:



This code is an implementation of a Random Forest Classifier model to predict customer churn in a Telco company. We use the pandas library to load the dataset from an CSV file and preprocess it. The `train_test_split` function from `sklearn.model_selection` is used to split the dataset into training and testing sets. The `RandomForestClassifier` class from `sklearn.ensemble` is used to initialize the model with 100 estimators. The model is then trained on the training set using the `fit` method. The `predict` method is used to make predictions on the testing set. Finally, the `accuracy_score`, `confusion_matrix`, and `classification_report` functions from `sklearn.metrics` are used to evaluate the model's performance. The confusion matrix is plotted using the `heatmap` function from `seaborn` and `matplotlib` libraries.


The dataset used contains information about customers who left within the last month. The dataset has 21 columns and 7,043 rows. The target variable is the Churn Label column, which indicates whether a customer has left or not. The features used in the model are all columns except for Churn Label, Churn Value, Churn Score, CLTV, and Churn Reason. The categorical variables in the dataset are converted to dummy/indicator variables using the `get_dummies` function from pandas.

The model's accuracy is 0.7877, or 79%. The confusion matrix shows the number of true positives, false positives, true negatives, and false negatives. The classification report shows the precision, recall, f1-score, and support for each class. The heatmap shows the confusion matrix in a graphical format.

The model's performance was improved by adding two more components to the code: feature importance plot and ROC curve.

The feature importance plot shows the relative importance of the top 10 features in the model. The ROC curve shows the trade-off between the true positive rate and the false positive rate for different threshold values. The AUC score is a measure of the model's performance, with a higher score indicating better performance. The AUC score for this model is 0.8536, or 85%.





Iris

Code explanation:

This code demonstrates a common approach to solve a multi-class classification problem using the Random Forest Classifier. It begins by loading the Iris dataset and preprocessing it.

Next, it splits the dataset into training and testing sets using the `train_test_split` function from `sklearn`. After that, it creates a Random Forest Classifier with 100 trees and trains it using the training data. Then, it makes predictions on the testing data and calculates the accuracy score by comparing the predictions with the actual values.

It also prints the confusion matrix, which is a table that describes the performance of the classifier by comparing predicted values with actual values.

The code further plots the confusion matrix and the feature importance, which represents the value of each feature in improving the classifier's performance.

Here is the breakdown of the code:

1. Load the Iris dataset using the pandas library.
2. Convert the 'species' column into numerical values using the `LabelEncoder` class from `sklearn.preprocessing`.
3. Split the dataset into training and testing sets using the `train_test_split` function from `sklearn.model_selection`.
4. Create a Random Forest Classifier with 100 trees using the `RandomForestClassifier` class from `sklearn.ensemble`.
5. Train the classifier using the training data.
6. Make predictions on the testing data.
7. Calculate the accuracy score of the classifier using the `accuracy_score` function from `sklearn.metrics`.
8. Print the confusion matrix and accuracy score.
9. Plot the confusion matrix using the seaborn library.
10. Plot the feature importance of the Random Forest Classifier.

